

DD2424 Deep Learning in Data Science

Assignment 3

Option 1: Batch Normalization

Ramona Häuselmann

May 6, 2021

1 Gradient Computation

I managed to successfully write the functions to correctly compute the gradient analytically. I tested my implementation by computing the gradients using `ComputeGradsNumSlow.m`. Then I compared my results with the numerical approaches by calculating the absolute difference of each gradient element as in equations below. Then I checked those against a threshold ($1e-5$). When using the reduced set with $n=2$ on a 3 layer network with $[50, 50]$ nodes in the hidden layers and using `ComputeGradsNumSlow.m` I get a maximum error of

- `diff_W1_max` = $2.2204e-11$
- `diff_b1_max` = $2.2204e-11$
- `diff_gamma1_max` = $2.2356e-11$
- `diff_beta1_max` = $2.2356e-11$
- `diff_W2_max` = $4.1234e-13$
- `diff_b2_max` = $2.2204e-16$
- `diff_gamma2_max` = $3.3339e-11$
- `diff_beta2_max` = $3.3339e-11$
- `diff_W3_max` = $2.3122e-11$
- `diff_b3_max` = $2.4201e-11$

These errors are small enough to conclude that my implementation works.

$$diff_W = abs(ngrad_W - grad_W) \quad (1)$$

$$diff_b = abs(ngrad_b - grad_b) \quad (2)$$

$$diff_gamma = abs(ngrad_gamma - grad_gamma) \quad (3)$$

$$diff_beta = abs(ngrad_beta - grad_beta) \quad (4)$$

2 3-layer Network

Training a 3-layer network with [50, 50] hidden nodes, $\lambda=0.005$, $\eta_{\min}=1e-5$, $\eta_{\max}=1e-1$, cycles=2, n_batch=100, n_s=5*45000/n_batch.

2.1 Without batch normalization

accuracy_validation = 54.1%

accuracy_test = 53.05%




Figure 1: Loss, 3-layer without batch normalization

2.2 With batch normalization

accuracy_validation = 55.1%

accuracy_test = 53.95%




Figure 2: Loss, 3-layer with batch normalization

3 9-layer Network

Training a 9-layer network with [50; 30; 20; 20; 10; 10; 10; 10] hidden nodes, $\lambda=0.005$, $\eta_{\min}=1e-5$, $\eta_{\max}=1e-1$, $\text{cycles}=2$, $n_{\text{batch}}=100$, $n_s=5 \cdot 45000/n_{\text{batch}}$.

3.1 Without batch normalization

$\text{accuracy_validation} = 47.1\%$

$\text{accuracy_test} = 45.18\%$




Figure 3: Loss, 9-layer without batch normalization

3.2 With batch normalization

$\text{accuracy_validation} = 52.3\%$

$\text{accuracy_test} = 51.6\%$




Figure 4: Loss, 9-layer with batch normalization

4 Optimizing 3-layer network

To optimize I first performed a coarse lambda search over a uniform grid between $\text{lambda}=1\text{e-}5$ and $\text{lambda}=1\text{e-}1$. I sampled 20 lambdas in the interval and trained for 2 cycles. This resulted in a best result (based on validation accuracy) at $\text{lambda}=0.00527$ and $\text{accuracy}=55.5\%$. Then I performed a finer, random search between $\text{lambda}=0.00001$ and $\text{lambda}=0.01$. I sampled 20 lambdas randomly and trained for 2 cycles. This resulted in a best result (based on validation accuracy) at $\text{lambda}=0.00564$ and $\text{accuracy}=55.6\%$. Then I trained the network with this lambda for 3 cycles, which resulted in a test accuracy of 54.27% and a validation accuracy of 56.1%.




Figure 5: Loss, optimized 3-layer with batch normalization




Figure 6: Accuracy, optimized 3-layer with batch normalization

5 Sensitivity to initialization

Instead of He initialization initialize the weights of the network with a normal distribution with the same sigma on each layer. Test this for 3 different sigmas.

5.1 sig=1e-1

5.1.1 Without batch normalization

accuracy_validation = 53.2%

accuracy_test = 53.24%




Figure 7: Loss, sig=1e-1, 3-layer without batch normalization

5.1.2 With batch normalization

accuracy_validation = 53.7%

accuracy_test = 53.49%




Figure 8: Loss, sig=1e-1, 3-layer with batch normalization

5.2 sig=1e-3

5.2.1 Without batch normalization

accuracy_validation = 50.6%

accuracy_test = 50.29%




Figure 9: Loss, sig=1e-3, 3-layer without batch normalization

5.2.2 With batch normalization

accuracy_validation = 54.9%

accuracy_test = 53.88%




Figure 10: Loss, sig=1e-3, 3-layer with batch normalization

5.3 sig=1e-4

5.3.1 Without batch normalization

accuracy_validation = 7.8%

accuracy_test = 10%




Figure 11: Loss, sig=1e-4, 3-layer without batch normalization

5.3.2 With batch normalization

accuracy_validation = 55.2%

accuracy_test = 53.44%




Figure 12: Loss, sig=1e-4, 3-layer with batch normalization

5.4 Conclusion

From the experiment we can see that batch normalization makes the training much more stable. The effect is best visible in the experiment with $\text{sig}=1\text{e-}4$. The initialization is very bad for training and without batch normalization the network achieves a very low test accuracy of only 10%. If we use batch normalization we achieve 53.44%, which is (almost) the same test accuracy as the optimized network which used a good initialization (Xavier).